

Any questions to exercises and homeworks from last time?

5 Sinusoidal Signals and the Fast Fourier Transform

- Spectrum Estimation using the FFT
- FFT Libraries in Python
- Partial Spectrum DFT Algorithms in Python
- Windowing
- Windowing using Python
- Generating a Signal with Multiple Sinusoids
- FFT-based Spectrum Estimation
- FFT-based Sinusoidal Parameter Estimation
 - Argmax Method and Zero-Padding
 - Polynomial Fit

The Fast Fourier Transform

- The FFT is an algorithm to compute the discrete time Fourier transform (DTFT) efficiently.
- The DTFT is the Fourier transform for equally spaced sampled complex data (in time and frequency domain)
- The DTFT is often used to find an approximation to the power spectral density

$$P(\psi) = \left| \frac{1}{N} \sum_{n=0}^{N-1} x[n] \exp(-j2\pi\psi n) \right|^2$$

at discrete sampling points k in the frequency domain

$$P[k] = \left| \frac{1}{N} \sum_{n=0}^{N-1} x[n] \exp\left(-j2\pi \frac{k n}{N}\right) \right|^2 \quad \text{for } k = 0, \dots, N-1$$

The Fast Fourier Transform in Python

- Implementation of the FFT in libraries may differ in pre-factor (1 , $\frac{1}{N}$, $\frac{1}{2\pi}$, or square-roots of them), exponent ($\pm 2\pi$ or ω), and start index (0 or $-\frac{N}{2}$), and the lacking of the absolute value.
- In Python, the FFT is usually calculated as

$$X[k] = \sum_{n=0}^{N-1} x[n] \exp\left(-2\pi j \frac{n k}{N}\right) \quad k = 0, \dots, N-1$$

which is just the sum-term of the power spectrum $P[k]$.

- Symmetric Fourier transform (start index at $-N/2$):

$$X_{\text{sym}} = X[k] \exp\left(-2\pi j \frac{k}{N} \frac{N-1}{2}\right)$$

- The FFT is not actually implemented in Python due to speed.
- FFT modules in Python just link to a compiled library.

FFT Libraries in Python

- Most common wrapper modules for the FFT are:

- `numpy.fft` <https://docs.scipy.org/doc/numpy-1.15.0/reference/routines.fft.html>
- `scipy.fftpack` <https://docs.scipy.org/doc/scipy/reference/fftpack.html>
- `pyfftw` <https://github.com/pyFFTW/pyFFTW>

- Functions provided by those modules have nearly the same arguments, thus modules are interchangeable.

- Speed of the FFT can vary heavily with

- utilized library (optimizations and multithreading),
- data type (single, or double precision), and
- number of samples (different algorithms perform differently)

- `scipy` and `numpy` work out of the box and try to use Intel's Math Kernel Library (MKL).

- FFTW is usually the fastest but needs extra steps on Windows, as Windows does not ship with a C compiler.

Goertzel-Algorithm

- To calculate the DFT for an arbitrary list of frequencies.
- Faster than FFT for low number of desired frequencies.
- No module in python but implementation is trivial.

Chirp-z Transform

- Also called chirp-transform algorithm (CTA), or zoom-FFT
- Based on Bluestein FFT (for arbitrary data length)
- CZT didn't make it into scipy/numpy due to numerical problems with corner cases.

<https://github.com/scipy/scipy/issues/4288>

- Available at <https://gist.github.com/ewmoore/6d321c1dbb5cea9bfaeb>

- There is no unwindowed FFT!

$$P_{\text{rect}}[k] = \left| \frac{1}{N} \sum_{n=0}^{N-1} x[n] \exp\left(-j2\pi \frac{nk}{N}\right) \right|^2$$

$$P_{\text{wind}}[k] = \left| \frac{1}{\sum_{n=0}^{N-1} a[n]} \sum_{n=0}^{N-1} a[n] x[n] \exp\left(-j2\pi \frac{nk}{N}\right) \right|^2$$

- Different windows produce different the PSF.
- Allows to trade beam-width for side-lobe level, e.g., increases dynamic range but reduces resolution.
- Note that the sum can still be evaluated using an FFT library!

Windowing using Python

- Three modules provide window functions:

- `numpy`
- `scipy.signal.windows`
- `spectrum.window`

- `numpy` provides some basic window functions with

`np.bartlett`, `np.blackman`, `np.hamming`, `np.hanning`, `np.kaiser`

```
import numpy as np
wind=np.hanning(128)
```

- `scipy.signal.windows` provide most commonly used window functions with the option for symmetric windows.

<https://docs.scipy.org/doc/scipy/reference/signal.windows.html>

```
import scipy.signal.windows as windows
wind=windows.nuttall(N, sym=True)
```


Windowing using Python

- `spectrum.window` provides many tools for spectrum estimation, like functions for generating and analysis of a wider range of window functions. <https://github.com/cokelaer/spectrum>

```
import spectrum
import matplotlib.pyplot as plt
from matplotlib2tikz import save as tikz_save





# generate the samples
wind_samples=spectrum.window.window_nuttall(128)

# generate window object
for name in ('hanning', 'nuttall'):
    wind=spectrum.Window(128,name)           # generate a window object
    wind_samples=wind.data                   # get the data

    plt.figure()
    wind.plot_time_freq()                   # plot the window function

    tikz_save('Window_'+name+'.tikz');
    plt.savefig('Window_'+name+'.png', dpi=150)
```

Windowing using Python



Sums of Sinusoidal Signals With Normalized Frequency

- Sum of M sinusoidal signal at sampling indices n :

$$s_{\text{real}}[n] = \sum_{m=0}^M A_m \sin(2\pi\psi_m n + \phi_m)$$


$$s_{\text{analytic}}[n] = \sum_{m=0}^M A_m \exp(j2\pi\psi_m n + j\phi_m)$$

- Normalized notations to become independent of sample rate.

- Here: $\psi_m = f_m / T_s$ with $\begin{array}{ll} \psi_m & \text{normalized frequency} \\ f_m & \text{frequency in Hz} \\ T_s & \text{sampling frequency in Hz} \end{array}$
- $\psi_m = 0$ is DC, $\psi_m = 0.5$ is half the sampling frequency (cf. Nyquist–Shannon sampling theorem)
- Note: $\sin x = \frac{1}{2j} (\exp(jx) - \exp(-jx))$


Exercise: Visualization of the FFT Spectrum

- ▶ Consider 32 samples of a complex signal taken at a sampling rate of 5 MHz with $A = 2$, $\psi = 0.14$, and $\phi = 0$.
- ▶ Estimate and plot the power spectrum using the FFT.
 - Plot in dB over MHz
 - What happens if ψ is a multiple of $1/N$, e.g., $\psi = 0.125$?
 - Hint: take a look at the keyword parameter `endpoint` if you use `linspace` to construct the frequency axis.



Exercise: Visualization of the FFT Spectrum

- ▶ Consider 32 samples of a complex signal taken at a sampling rate of 5 MHz with $A = 2$, $\psi = 0.14$, and $\phi = 0$.
- ▶ Estimate and plot the power spectrum using the FFT.
 - Plot in dB over MHz
 - What happens if ψ is a multiple of $1/N$, e.g. $\psi = 0.125$?
 - Hint: take a look at the keyword parameter `endpoint` if you use `linspace` to construct the frequency axis.



Homework: Visualization of the FFT Spectrum

- ▶ Consider 128 samples of a real signal taken at a sampling rate of 2 MHz, containing three sinusoids with



```
Am    = [1.0, 0.5, 1e-4, 0.001];  
psim  = [0.1, 0.1+2.5/N, 0.17, 0.21];  
phim  = [0.0, 0.0, 2.0, 0.0];
```

- ▶ Estimate the power spectrum by using the FFT.
 - Plot it in dB (normalized to $A_m[0]$) over frequency in MHz.
 - Use `xlim` and `ylim` to zoom the plot to the interesting region 0.1 to 0.5 MHz and -100 to 10 dB.
 - Mark the magnitudes and frequencies in the spectrum plot.
- ▶ Plot the power spectrum for three different window functions: rect, hanning and nuttall.
- ▶ The ideal Fourier transform would produce 4 peaks. Explain:
 - Why (in terms of shape of the main- and sidelobes) peaks are not visible/distinct for certain window functions?
 - Why do the tips of the DFT/FFT peaks not meet the sinusoidal parameters?

Sinusoidal Parameter Estimation Using the FFT

- Idea: Location of the maximum in the power spectrum corresponds to the parameters of the strongest sinusoid.
- Zero padding
 - Padding the signal with zeros prior the FFT produces a finer spectral grid.
 - The function call to FFT has a parameter for zero padding. NO manual concatenate needed.
 - The pulse-spread-function does not change!
 - Thus, discretization errors are reduced, resolution is not increased, however memory consumption is increased.
- AWGN introduces additional errors. Similar to the AWGN peak detection from the correlation exercise.

Zero-Padding Example



- The frequency discretization of the DTFT is—one of many—sources for an inaccurate frequency estimation using the argmax method.
- The main lobe of the PSD (squared magnitude spectrum) is well approximated by a 2nd-order polynomial for most window functions.

$$P(\psi) = a_0 + a_1\psi + a_2\psi^2$$

- The location of the maximum of the 2nd polynomial can be found analytical without discretization.

$$0 = \left. \frac{dP(\psi)}{d\psi} \right|_{\psi=\psi_{\max}} \Rightarrow \psi_{\max} = -\frac{a_1}{2a_2} \quad F_{\max} = a_0 - \frac{a_1^2}{4a_2}$$

- No parabola/linear fit for phase due to wrapping on 2π . Use DFT/goertzel ones ψ_{\max} found.

Calculate Parameters of the Polynomial

■ Tricks for derivation:

- center the system of coordinates around 0 frequency
- FFT bins are all equally spaced by d_f , i.e. $\psi \in \{0, -d_f, d_f\}$

$$\begin{pmatrix} P_L \\ P_C \\ P_R \end{pmatrix} = \begin{pmatrix} 1 & -d_f & d_f^2 \\ 1 & 0 & 0 \\ 1 & d_f & d_f^2 \end{pmatrix} \begin{pmatrix} a_0 \\ a_1 \\ a_2 \end{pmatrix}$$

Calculate Parameters of the Polynomial

■ Tricks for derivation:

- center the system of coordinates around 0 frequency
- FFT bins are all equally spaced by d_f , i.e. $\psi \in \{0, -d_f, d_f\}$

$$\begin{pmatrix} P_L \\ P_C \\ P_R \end{pmatrix} = \begin{pmatrix} 1 & -d_f & d_f^2 \\ 1 & 0 & 0 \\ 1 & d_f & d_f^2 \end{pmatrix} \begin{pmatrix} a_0 \\ a_1 \\ a_2 \end{pmatrix}$$

■ Polynomial Parameters:

$$a_0 = P_C \quad a_1 = \frac{P_R - P_L}{2d_f} \quad a_2 = \frac{P_L - 2P_C + P_R}{2d_f^2}$$

Calculate Parameters of the Polynomial

■ Tricks for derivation:

- center the system of coordinates around 0 frequency
- FFT bins are all equally spaced by d_f , i.e. $\psi \in \{0, -d_f, d_f\}$

$$\begin{pmatrix} P_L \\ P_C \\ P_R \end{pmatrix} = \begin{pmatrix} 1 & -d_f & d_f^2 \\ 1 & 0 & 0 \\ 1 & d_f & d_f^2 \end{pmatrix} \begin{pmatrix} a_0 \\ a_1 \\ a_2 \end{pmatrix}$$

■ Polynomial Parameters:

$$a_0 = P_C \quad a_1 = \frac{P_R - P_L}{2d_f} \quad a_2 = \frac{P_L - 2P_C + P_R}{2d_f^2}$$

■ Result:

$$\psi_{\text{offset}} = -\frac{a_1}{2a_2} = \frac{d_f(P_R - P_L)}{2(P_L - 2P_C + P_R)}$$
$$F_{\text{max}} = \frac{-2P_R(4P_C + P_L) + (P_L - 4P_C)^2 + P_R^2}{16P_C - 8(P_L + P_R)}$$

Calculate Parameters of the Polynomial

■ Tricks for derivation:

- center the system of coordinates around 0 frequency
- FFT bins are all equally spaced by d_f , i.e. $\psi \in \{0, -d_f, d_f\}$

$$\begin{pmatrix} P_L \\ P_C \\ P_R \end{pmatrix} = \begin{pmatrix} 1 & -d_f & d_f^2 \\ 1 & 0 & 0 \\ 1 & d_f & d_f^2 \end{pmatrix} \begin{pmatrix} a_0 \\ a_1 \\ a_2 \end{pmatrix}$$

■ Polynomial Parameters:

$$a_0 = P_C \quad a_1 = \frac{P_R - P_L}{2d_f} \quad a_2 = \frac{P_L - 2P_C + P_R}{2d_f^2}$$

■ Result:

$$\psi_{\text{offset}} = -\frac{a_1}{2a_2} = \frac{d_f(P_R - P_L)}{2(P_L - 2P_C + P_R)}$$
$$F_{\text{max}} = \frac{-2P_R(4P_C + P_L) + (P_L - 4P_C)^2 + P_R^2}{16P_C - 8(P_L + P_R)}$$

■ Zero-padding is still required but to a less extent.

Exercise: Calculate Polynomial Fit

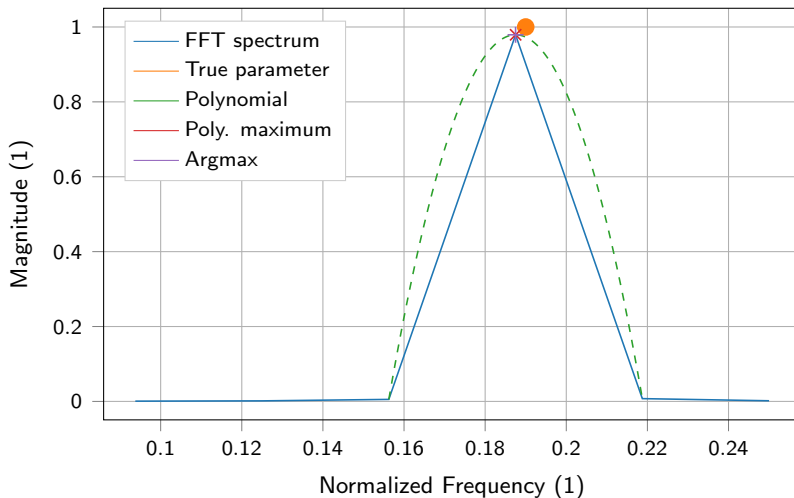
```
def calc_polyfit_correction(spectrum_excert, df):  
    """  
    Parameters  
    -----  
    spectrum_excert: list of three floats  
        left, center and right magnitude value  
  
    df: float  
        frequency difference of two consecutive FFT-bins  
  
    Returns  
    -----  
    offset_frequency: float  
        Offset towards the location of Fc  
  
    Fmax: magnitude at the maximum  
  
    a: list of three floats  
        Parameters of the polynomial.  
  
    """
```

Exercise: Test Polynomial Fit

- ▶ Generate and plot the FFT based spectrum estimate for a sinusoid `A=[1]; psi=[0.19]; phi=[0.85]; N=32; Z=4*N`.
- ▶ Apply the method `calc_polynomial_fit` from the exercise before.
- ▶ Annotate the plot with
 - a marker for the sinusoidal parameters,
 - the maximum using `argmax`,
 - the maximum obtained by polynomial fit, and
 - the polynomial in a region of ± 1 FFT bin around the maximum.
- ▶ Hints:
 - `numpy.polyval` can help drawing the polynomial curve
 - The FFT spectrum is periodic, thus getting the left and right value might need index wrapping.
 - If `argmax` produced index 0, negative indexing allows to get the left value with `index-1`.
 - If `argmax` produced index $Z - 1$, obtaining the right value might produce `index is out of bounds` error. Check for this case or use the modulo operation `%` to avoid it.

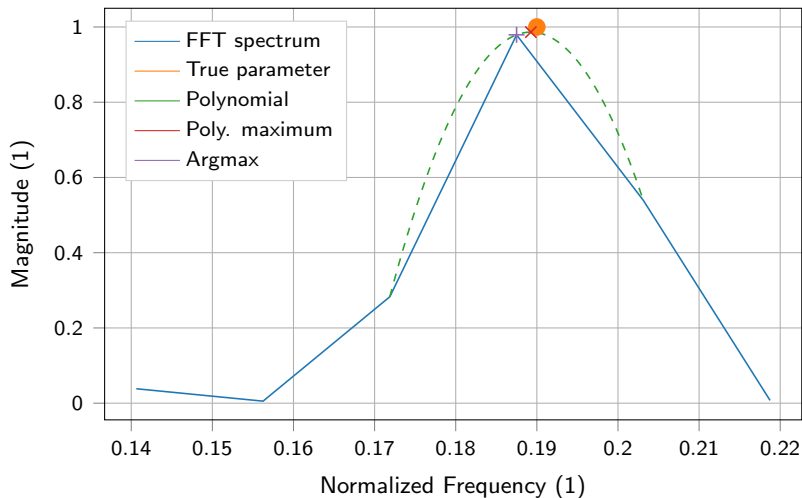
Solution to the Polyfit Exercise

$N=32$, $Z=32$, $\psi_{\text{err,argmax}}/\psi_{\text{err,polyfit}}$: $2.50\text{e-}03/2.48\text{e-}03$




Solution to the Polyfit Exercise

$N=32$, $Z=64$, $\psi_{\text{err,argmax}}/\psi_{\text{err,polyfit}}$: $2.50\text{e-}03/7.38\text{e-}04$



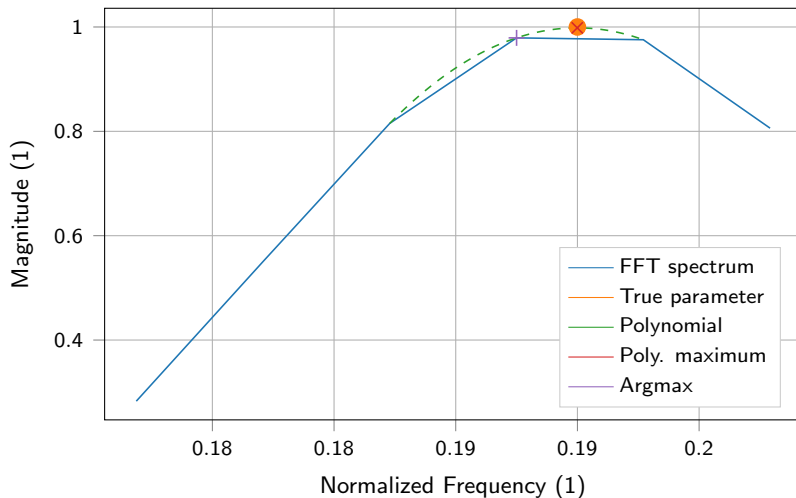
Solution to the Polyfit Exercise

$N=32$, $Z=128$, $\psi_{\text{err,argmax}}/\psi_{\text{err,polyfit}}: 2.50\text{e-}03/1.25\text{e-}04$



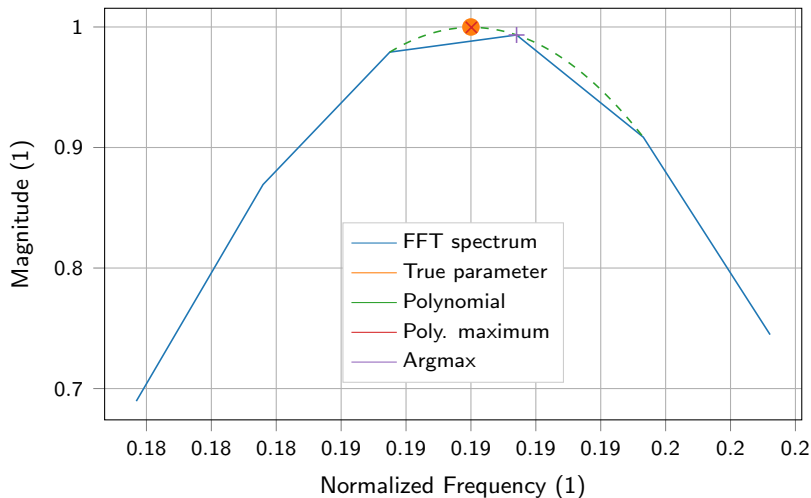
Solution to the Polyfit Exercise

$N=32$, $Z=192$, $\psi_{\text{err,argmax}}/\psi_{\text{err,polyfit}}$: $2.50\text{e-}03/7.40\text{e-}06$




Solution to the Polyfit Exercise

$N=32$, $Z=256$, $\psi_{\text{err,argmax}}/\psi_{\text{err,polyfit}}$: $1.41\text{e-}03/1.40\text{e-}05$




Solution to the Polyfit Exercise

$N=32$, $Z=512$, $\psi_{\text{err,argmax}}/\psi_{\text{err,polyfit}}: 5.47\text{e-}04/1.93\text{e-}06$



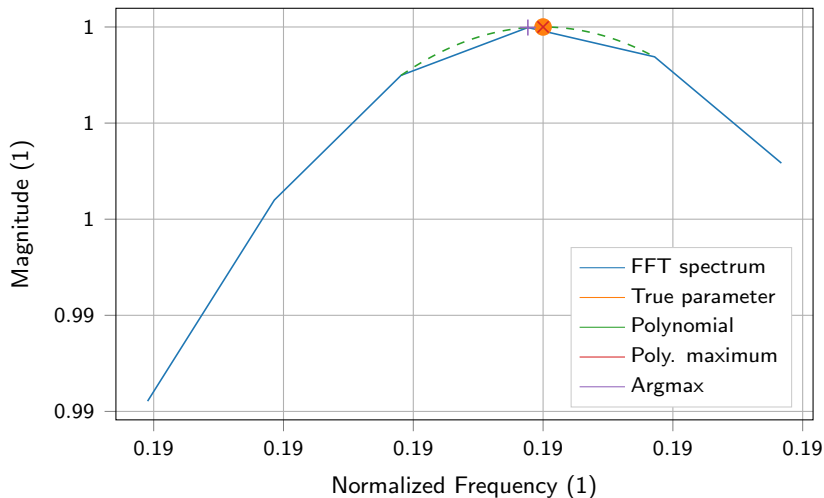
Solution to the Polyfit Exercise

$N=32$, $Z=1024$, $\psi_{\text{err,argmax}}/\psi_{\text{err,polyfit}}$: $4.30\text{e-}04/1.25\text{e-}07$



Solution to the Polyfit Exercise


$N=32$, $Z=2048$, $\psi_{\text{err,argmax}}/\psi_{\text{err,polyfit}}$: $5.86\text{e-}05/1.77\text{e-}08$



Homework: Phase estimate

- ▶ Implement either the DFT sum or the Goertzel algorithm to be able to estimate a certain spectral bin.
- ▶ Use the DFT sum or the Goertzel algorithm to estimate the phase at the frequency obtained by the parabola fit.
- ▶ Plot the phase spectrum and mark the true and the estimated phase.

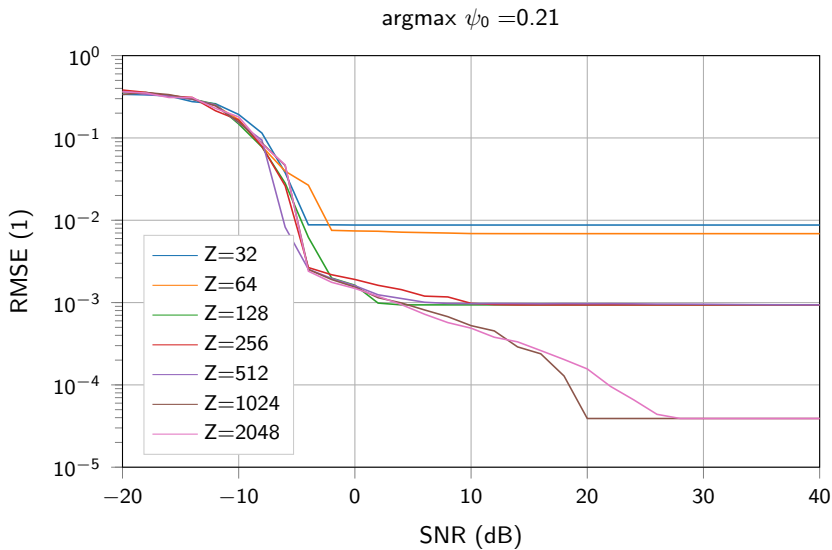
$N=32$, $Z=128$, $\phi_{\text{err,argmax}}/\phi_{\text{err,polyfit}}$: $2.43\text{e-}01/1.22\text{e-}02$



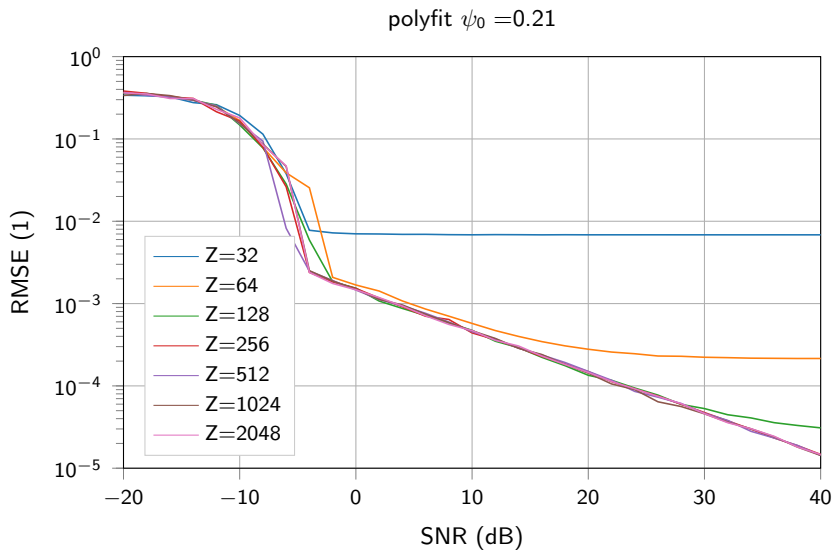
Homework: RMSE of Argmax, and Polyfit

- ▶ Consider a signal $s_{RX} = s_{IF} + n$ of length $N = 32$, where s_{IF} is a complex cisoid with $A = 1$, $\psi = 0.21$, $\phi = 0$ and n is zero-mean complex white Gaussian noise.
- ▶ Plot the RMSE (in logarithmic scale) of frequency estimation over SNR (-20 to 40 dB in steps of 2 dB) using the argmax method and the polyfit method for $Z_f \in \{1, 2, 4, 8, 16, 32, 64\}$.
- ▶ Hints:
 - $SNR = \frac{A^2}{2\sigma^2}$
 - Z is the zero-padding factor, i.e. a signal with N should produce $Z = Z_f N$ FFT-bins.
 - Use at least 1000 different noise realizations for each SNR to get smooth curves for the RMSE.

Expected Plots for Homework



Expected Plots for Homework



Homework: RMSE of Argmax, and Polyfit, cntd.

- ▶ Why is a higher zero-padding factor more beneficial for the argmax method than for polyfit?
- ▶ Why does the polyfit method not decrease the RMSE (compared to argmax) for $Z < 4$?
- ▶ Why can the argmax method show the same RMSE (at high SNR) for different values of Z ?
- ▶ Why does the achievable RMSE change when using the argmax method (consider $\psi \in \{0.21, 0.22, 0.23\}$)?